

All-flash VAST storage, delivered as native Kubernetes storage

ISC stands up VAST Data on AWS and wires it into SpectroCloud Palette, so every cluster gets file, block, and object storage from one disaggregated all-flash platform — with per-tenant isolation and QoS built in.

Deck-ready · version-agnostic

Drops into v4b_naming / v5

One slide, two if split

THE PITCH · 30 SECONDS

VAST Data is an all-flash, exabyte-scale platform built for AI/ML, analytics, and HPC. Delivered on AWS as **VAST-on-Cloud** and integrated into **Palette**, a single VAST cluster serves **NFS file, NVMe/TCP block, and S3 object** to any number of Kubernetes clusters — each carved into its own isolated tenant. One high-performance data platform under every workload, instead of three storage products stitched together.

WHAT IT DELIVERS · THREE PROTOCOLS, ONE PLATFORM

PROTOCOL	KUBERNETES SURFACE	ACCESS	BEST FOR
NFS filesystem csi.vastdata.com · vastdata-filesystem	CSI persistent volumes	RWX · shared	AI/ML training data, shared datasets, notebooks
NVMe/TCP block block.csi.vastdata.com · vastdata-block	CSI persistent volumes	RWO	Databases, low-latency stateful apps
S3 object objectstorage.k8s.io · COSI	Bucket-per-workload	object	Data lakes, model artifacts, RAG corpora

01 One platform, three protocols File, block, and object from the same all-flash cluster. No tiering, no separate object store, no separate block array.	02 Multi-tenancy that maps to how customers segment Each Palette cluster gets an isolated VAST tenant — dedicated VIP pool, view policy, storage path, and QoS tier (bronze / silver / gold) capping bandwidth and IOPS.
03 Lifecycle-managed, not bolted on The VAST drivers ship as Palette cluster-profile packs — storage is versioned and deployed with the cluster, repeatable across every environment.	04 Built for the AI workload All-flash throughput and NVMe/TCP latency are what training and inference pipelines need; VAST's data reduction keeps the economics sane at scale.

TECH KEYWORDS TO ADD

VAST Data	VAST-on-Cloud	DASE	all-flash	NVMe/TCP	CSI	COSI
S3 object storage	RWX / RW0 PVs	storage multi-tenancy	QoS (bandwidth/IOPS)			
tenant isolation	Kubernetes-native storage	SpectroCloud Palette				
cluster profiles	AI/ML data pipeline	RAG	inference	data lake		

PROOF POINTS · VALIDATED, NOT ASPIRATIONAL

✓ NFS RWX PASS Provisioning plus read/write verified end-to-end on a live Palette cluster against an isolated VAST tenant.	✓ Block & object INTEGRATED NVMe/TCP block and COSI/S3 drivers running; block volumes provision and attach over the tenant data path.
✓ Tenant isolation PROVEN Distinct VIP pools, views, and QoS per tenant across separate Palette clusters, gated by client-network scoping.	✓ Codified stack REPEATABLE Terraform plus Palette packs — demonstrable and repeatable in a customer account, not a one-off lab.

AI/ML platform teams, data-intensive SaaS, HPC and research, and any Kubernetes shop running three storage products where one would do. Lead use cases: AI training and inference data platforms, RAG and vector pipelines, multi-tenant internal developer platforms.

AI platform teams

data-intensive SaaS

HPC / research

internal dev platforms

“One all-flash platform delivering file, block, and object to every Kubernetes cluster — each team isolated in its own tenant with its own performance tier, deployed and versioned with the cluster itself.”